

# Modeling Run Performance

Joe Martin

10/27/2021

## Garmin Data Modeling

The two most obvious primary target variables are average speed (avg\_spd) in miles per hour, and average pace (avg\_pace\_sec) in seconds. A higher average speed and a lower average pace are the desired outcome when measuring performance over time. Reviewing the results of the two preliminary linear regression models, the more desirable variable is average pace, as it has stronger relationships with other variables.

```
# Create preliminary model
```

```
prelim_spd <- lm(avg_spd ~ ., df)
summary(prelim_spd)
```

```
##
## Call:
## lm(formula = avg_spd ~ ., data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.187959 -0.029389 -0.000816  0.028251  0.219902
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -6.402e+00  5.019e-01 -12.755 < 2e-16 ***
## distance       1.031e-02  7.272e-03   1.417 0.157270
## avg_hr         4.347e-03  1.026e-03   4.236 2.84e-05 ***
## max_hr        -1.052e-03  6.662e-04  -1.579 0.115059
## avg_run_cadence 4.733e-02  1.259e-03  37.584 < 2e-16 ***
## max_run_cadence 1.806e-04  2.388e-04   0.756 0.449921
## total_ascent   -5.052e-05  7.741e-05  -0.653 0.514341
## total_decent    6.208e-05  7.328e-05   0.847 0.397386
## avg_stride      5.142e+00  1.294e-01  39.720 < 2e-16 ***
## min_elevation   3.616e-04  1.058e-04   3.418 0.000697 ***
## max_elevation  -2.596e-05  1.096e-04  -0.237 0.812852
## avg_pace_sec    -1.011e-03  3.465e-04  -2.916 0.003745 **
## best_pace_sec   -9.307e-05  1.039e-04  -0.896 0.370737
## 'sweat_loss(ml)' -4.945e-05  1.336e-04  -0.370 0.711397
## aerobic_TE      -8.122e-02  1.663e-02  -4.885 1.51e-06 ***
## aerobic_fctImpacting -5.322e-03  9.255e-03  -0.575 0.565576
## aerobic_fctMaintaining 2.091e-02  1.771e-02   1.181 0.238502
## aerobic_fctOverreaching 4.604e-02  1.427e-02   3.225 0.001365 **
## anaerobic_value  1.303e-02  1.097e-02   1.187 0.235838
```

```
## anaerobic_fctMaintaining -4.067e-03 1.741e-02 -0.234 0.815458
## anaerobic_fctNo Benefit -4.778e-02 3.440e-02 -1.389 0.165630
## anaerobic_fctSome Benefit -6.501e-02 2.583e-02 -2.517 0.012246 *
## max_spd -3.014e-03 2.567e-03 -1.174 0.241070
## short_distanceY 1.336e-02 1.906e-02 0.701 0.483817
## middle_distanceY 1.843e-02 1.414e-02 1.303 0.193186
## long_distanceY NA NA NA NA
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.04525 on 392 degrees of freedom
## (11 observations deleted due to missingness)
## Multiple R-squared: 0.9978, Adjusted R-squared: 0.9976
## F-statistic: 7263 on 24 and 392 DF, p-value: < 2.2e-16
```

```
prelim_pace <- lm(avg_pace_sec ~ ., df)
summary(prelim_pace)
```

```
##
## Call:
## lm(formula = avg_pace_sec ~ ., data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31.013  -3.313  -0.427   3.011  47.420
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.185e+03  6.190e+01  19.143 < 2e-16 ***
## distance       -4.443e+00  1.027e+00  -4.326 1.93e-05 ***
## avg_hr         3.786e-01  1.501e-01   2.522 0.01206 *
## max_hr        -8.380e-03  9.638e-02  -0.087 0.93075
## avg_run_cadence -1.787e+00  3.790e-01  -4.715 3.37e-06 ***
## max_run_cadence  5.933e-03  3.445e-02   0.172 0.86337
## total_ascent    -9.885e-03  1.116e-02  -0.886 0.37613
## total_decent     3.999e-03  1.057e-02   0.378 0.70547
## avg_stride     -2.333e+02  4.015e+01  -5.811 1.29e-08 ***
## min_elevation   -2.627e-02  1.542e-02  -1.703 0.08930 .
## max_elevation    2.642e-02  1.575e-02   1.678 0.09417 .
## best_pace_sec    2.566e-02  1.494e-02   1.718 0.08661 .
## 'sweat_loss(ml)'  9.787e-02  1.862e-02   5.257 2.41e-07 ***
## aerobic_TE      -8.762e+00  2.429e+00  -3.606 0.00035 ***
## aerobic_fctImpacting -4.257e+00  1.318e+00  -3.231 0.00134 **
## aerobic_fctMaintaining  5.590e+00  2.542e+00   2.199 0.02849 *
## aerobic_fctOverreaching  5.209e+00  2.069e+00   2.518 0.01220 *
## anaerobic_value  -1.135e+00  1.584e+00  -0.716 0.47435
## anaerobic_fctMaintaining  2.063e+00  2.509e+00   0.822 0.41144
## anaerobic_fctNo Benefit  1.926e-01  4.972e+00   0.039 0.96913
## anaerobic_fctSome Benefit  1.513e+00  3.754e+00   0.403 0.68709
## avg_spd        -2.101e+01  7.205e+00  -2.916 0.00374 **
## max_spd         6.393e-02  3.708e-01   0.172 0.86321
## short_distanceY  -5.004e-01  2.750e+00  -0.182 0.85569
## middle_distanceY -1.782e-01  2.044e+00  -0.087 0.93054
## long_distanceY    NA         NA         NA         NA
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.525 on 392 degrees of freedom
## (11 observations deleted due to missingness)
## Multiple R-squared:  0.9901, Adjusted R-squared:  0.9895
## F-statistic: 1641 on 24 and 392 DF,  p-value: < 2.2e-16
```

The ultimate goal of this model is to utilize data leading up to a performance event. As many races take place on Sunday and the typical long-distance run in this data set takes place on Sunday, the final linear regression model will begin with predicting Sunday performance.

To being predicting run performance, an initial linear regression model will be built below using all available data. Based on the preliminary linear regression above, an aerobic training effect that has a high impact (value between 4 and 4.9) is strongly related to average pace. This variable will be the target variable in the logistic regression that follows.

```
set.seed(456)
# Split data into training and testing sets
df_split <- initial_split(df, prop = 3/4)

train_df <- training(df_split)
test_df <- testing(df_split)

# Create recipe
pace_rec <- recipe(avg_pace_sec ~ ., data = train_df)

summary(pace_rec)
```

```
## # A tibble: 22 x 4
##   variable      type    role    source
##   <chr>        <chr>  <chr>  <chr>
## 1 distance      numeric predictor original
## 2 avg_hr        numeric predictor original
## 3 max_hr        numeric predictor original
## 4 avg_run_cadence numeric predictor original
## 5 max_run_cadence numeric predictor original
## 6 total_ascent  numeric predictor original
## 7 total_decent  numeric predictor original
## 8 avg_stride    numeric predictor original
## 9 min_elevation numeric predictor original
## 10 max_elevation numeric predictor original
## # ... with 12 more rows
```

```
lm_pace <- linear_reg() %>%
  set_engine("lm")

pace_wflow <- workflow()%>%
  add_model(lm_pace) %>%
  add_recipe(pace_rec)

pace_fit <- pace_wflow %>%
  fit(data = train_df)
```

```
tidy(pace_fit)
```

```
## # A tibble: 26 x 5
##   term                estimate std.error statistic  p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)      1166.        73.6      15.8 4.97e-41
## 2 distance         -3.81         1.22     -3.12 2.00e- 3
## 3 avg_hr            0.400        0.181      2.21 2.80e- 2
## 4 max_hr           -0.104        0.119     -0.880 3.80e- 1
## 5 avg_run_cadence  -1.70         0.456     -3.73 2.27e- 4
## 6 max_run_cadence   0.0245       0.0436      0.562 5.75e- 1
## 7 total_ascent     -0.00996      0.0139     -0.714 4.76e- 1
## 8 total_decent      0.00380      0.0130      0.293 7.70e- 1
## 9 avg_stride       -233.         48.9      -4.76 3.09e- 6
## 10 min_elevation   -0.0261       0.0193     -1.35 1.77e- 1
## # ... with 16 more rows
```

```
predict(pace_fit, test_df)
```

```
## Warning in predict.lm(object = object$fit, newdata = new_data, type =
## "response"): prediction from a rank-deficient fit may be misleading
```

```
## # A tibble: 107 x 1
##   .pred
##   <dbl>
## 1 450.
## 2 448.
## 3 436.
## 4 440.
## 5 397.
## 6 414.
## 7 427.
## 8 435.
## 9 433.
## 10 448.
## # ... with 97 more rows
```

```
pace_aug <- augment(pace_fit, test_df)
```

```
## Warning in predict.lm(object = object$fit, newdata = new_data, type =
## "response"): prediction from a rank-deficient fit may be misleading
```

```
pace_aug %>% select(avg_pace_sec, .pred)
```

```
## # A tibble: 107 x 2
##   avg_pace_sec .pred
##   <dbl> <dbl>
## 1      447 450.
## 2      449 448.
## 3      432 436.
```

```
## 4      438 440.
## 5      391 397.
## 6      414 414.
## 7      419 427.
## 8      432 435.
## 9      430 433.
## 10     444 448.
## # ... with 97 more rows
```

The R Mean-Squared Error for this model is 5.41. In other words, this model can predict average pace within 5.41 seconds.

```
pace_error <- pace_aug %>%
  rmse(truth = avg_pace_sec, .pred)

pace_error
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      5.41
```

The most significant variables (based on p-value) are average heart rate, average cadence, average stride, and aerobic training effect. The binary variable `aerobic_fct_Impacting` had a good p-value, as well, but that value is related to aerobic training effect, so it is left out of this analysis. As an attempt to improve the quality of the model, only the variables with the highest p-values will be included in this analysis.

```
set.seed(456)
# Split data into training and testing sets
df_split <- initial_split(df, prop = 3/4)

train_df <- training(df_split)
test_df <- testing(df_split)

# Create recipe
pace_rec_2 <- recipe(avg_pace_sec ~ avg_hr + avg_run_cadence + avg_stride + aerobic_TE, data = train_df)

summary(pace_rec_2)
```

```
## # A tibble: 5 x 4
##   variable      type    role    source
##   <chr>        <chr>  <chr>  <chr>
## 1 avg_hr      numeric predictor original
## 2 avg_run_cadence numeric predictor original
## 3 avg_stride   numeric predictor original
## 4 aerobic_TE   numeric predictor original
## 5 avg_pace_sec numeric outcome  original
```

```
lreg <- linear_reg() %>%
  set_engine("lm")

pace_wflow_2 <- workflow()%>%
```

```

add_model(lreg) %>%
add_recipe(pace_rec_2)

pace_fit_2 <- pace_wflow_2 %>%
  fit(data = train_df)

tidy(pace_fit_2)

```

```

## # A tibble: 5 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)    1460.      20.4      71.6 2.32e-197
## 2 avg_hr          0.215     0.141      1.53 1.27e- 1
## 3 avg_run_cadence -3.15     0.147     -21.5 1.01e- 63
## 4 avg_stride     -379.      7.16     -53.0 2.49e-159
## 5 aerobic_TE      -7.35     0.992     -7.41 1.16e- 12

```

```

predict(pace_fit_2, test_df)

```

```

## # A tibble: 107 x 1
##   .pred
##   <dbl>
## 1 452.
## 2 450.
## 3 444.
## 4 441.
## 5 388.
## 6 414.
## 7 430.
## 8 430.
## 9 442.
## 10 447.
## # ... with 97 more rows

```

```

pace_aug_2 <- augment(pace_fit_2, test_df)

pace_aug_2 %>% select(avg_pace_sec, .pred)

```

```

## # A tibble: 107 x 2
##   avg_pace_sec .pred
##   <dbl> <dbl>
## 1      447 452.
## 2      449 450.
## 3      432 444.
## 4      438 441.
## 5      391 388.
## 6      414 414.
## 7      419 430.
## 8      432 430.
## 9      430 442.
## 10     444 447.
## # ... with 97 more rows

```

Reviewing the results, the quality of the model decreased slightly. However, it seems that average pace will be a good target variable in exploring performance improvements.

```
pace_error_2 <- pace_aug_2 %>%  
  rmse(truth = avg_pace_sec, .pred)  
  
pace_error_2
```

```
## # A tibble: 1 x 3  
##   .metric .estimator .estimate  
##   <chr>   <chr>         <dbl>  
## 1 rmse    standard         6.50
```

## Logistic Regression

### All Variables

This first logistic regression is meant to predict whether an activity highly impacts aerobic training. This variable is relevant because it is the highest measure for aerobic conditioning without being over-reaching. In this analysis, calories and other variables related to aerobic training effect were removed.

```
set.seed(456)  
  
df2 <- df %>% mutate(high_impact = ifelse(aerobic_fct == "Highly Impacting", 1,0))  
df2$high_impact <- factor(df2$high_impact)  
  
# Split data into training and testing sets  
df2_split <- initial_split(df2, prop = 3/4)  
  
train_df2 <- training(df2_split)  
test_df2 <- testing(df2_split)  
  
# Create recipe. Use all variables except aerobic_TE and related  
aerobic_rec <- recipe(high_impact ~ short_distance + middle_distance + long_distance +  
  max_spd + avg_spd + anaerobic_value + `sweat_loss(ml)` + best_pace_sec +  
  avg_pace_sec + max_elevation + min_elevation + avg_stride +  
  total_decent + total_ascent + max_run_cadence + avg_run_cadence +  
  max_hr + avg_hr + distance, data = train_df2)  
  
summary(aerobic_rec)
```

```
## # A tibble: 20 x 4  
##   variable      type    role    source  
##   <chr>         <chr>  <chr>   <chr>  
## 1 short_distance nominal predictor original  
## 2 middle_distance nominal predictor original  
## 3 long_distance  nominal predictor original  
## 4 max_spd        numeric predictor original  
## 5 avg_spd        numeric predictor original  
## 6 anaerobic_value numeric predictor original  
## 7 sweat_loss(ml) numeric predictor original  
## 8 best_pace_sec  numeric predictor original
```

```
## 9 avg_pace_sec      numeric predictor original
## 10 max_elevation    numeric predictor original
## 11 min_elevation    numeric predictor original
## 12 avg_stride       numeric predictor original
## 13 total_decent     numeric predictor original
## 14 total_ascent     numeric predictor original
## 15 max_run_cadence  numeric predictor original
## 16 avg_run_cadence  numeric predictor original
## 17 max_hr           numeric predictor original
## 18 avg_hr           numeric predictor original
## 19 distance         numeric predictor original
## 20 high_impact      nominal outcome original
```

```
log_reg <- logistic_reg() %>%
  set_engine("glm")

aero_wkfl <- workflow()%>%
  add_model(log_reg) %>%
  add_recipe(aerobic_rec)

aero_fit <- aero_wkfl %>%
  fit(data = train_df2)

tidy(aero_fit)
```

```
## # A tibble: 20 x 5
##   term                estimate std.error statistic p.value
##   <chr>                <dbl>     <dbl>     <dbl>   <dbl>
## 1 (Intercept)        -111.      107.      -1.04    0.297
## 2 short_distanceY      0.114      1.34       0.0851  0.932
## 3 middle_distanceY     1.89       0.900      2.09    0.0363
## 4 long_distanceY       NA         NA         NA       NA
## 5 max_spd             -0.811     1.05      -0.775   0.439
## 6 avg_spd             -17.1     5.36      -3.19   0.00144
## 7 anaerobic_value     -0.670     0.425     -1.58    0.115
## 8 'sweat_loss(ml)'    -0.00296   0.0227    -0.130   0.896
## 9 best_pace_sec       -0.0278    0.0347    -0.800   0.423
## 10 avg_pace_sec        -0.0252    0.0851    -0.297   0.767
## 11 max_elevation       -0.0283    0.0118    -2.40    0.0165
## 12 min_elevation       0.0288    0.0102     2.81    0.00495
## 13 avg_stride          88.9     32.8       2.71    0.00673
## 14 total_decent        0.00536   0.00665    0.806    0.420
## 15 total_ascent        0.00340   0.00672    0.506    0.613
## 16 max_run_cadence     0.00974   0.0226    0.432    0.666
## 17 avg_run_cadence     0.747     0.298     2.50    0.0123
## 18 max_hr              0.0856    0.0742     1.15    0.249
## 19 avg_hr              0.115     0.0643     1.79    0.0734
## 20 distance           0.279     1.11      0.252    0.801
```

```
predict(aero_fit, test_df2)
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```



```
## # A tibble: 107 x 1
##   .pred_class
##   <fct>
## 1 0
## 2 0
## 3 0
## 4 0
## 5 0
## 6 0
## 7 0
## 8 1
## 9 0
## 10 1
## # ... with 97 more rows
```

```
aero_aug <- augment(aero_fit, test_df2)
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```

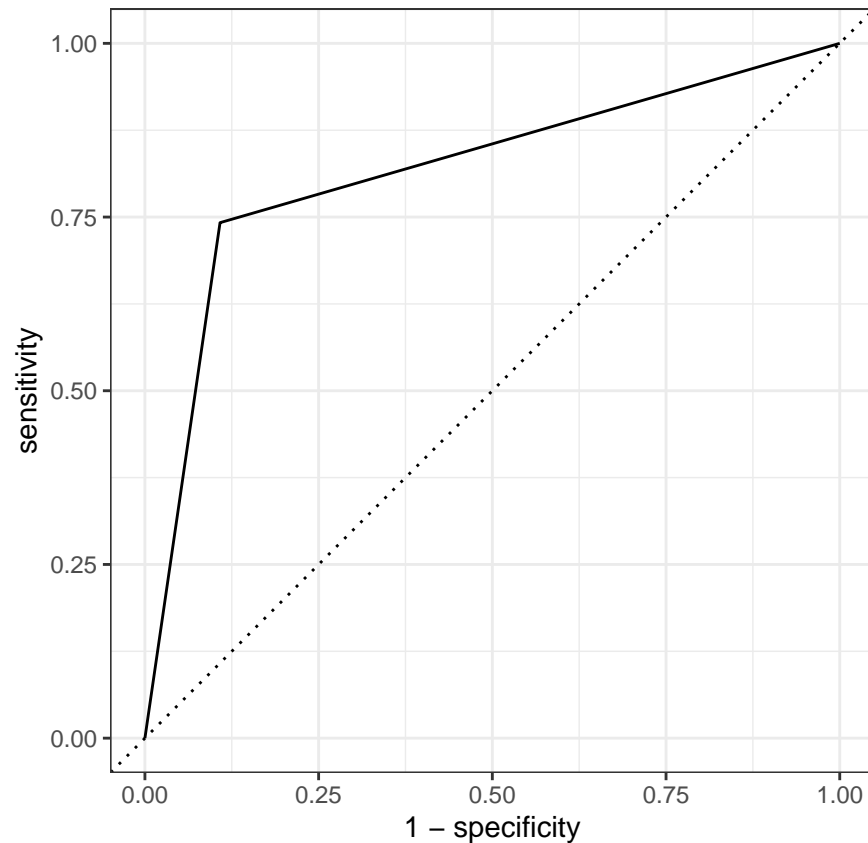
```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```

```
aero_aug %>% select(high_impact, .pred_class)
```

```
## # A tibble: 107 x 2
##   high_impact .pred_class
##   <fct>      <fct>
## 1 0          0
## 2 0          0
## 3 0          0
## 4 0          0
## 5 0          0
## 6 0          0
## 7 0          0
## 8 1          1
## 9 0          0
## 10 1         1
## # ... with 97 more rows
```

```
aero_aug$.pred_class <- as.character(aero_aug$.pred_class)
aero_aug$.pred_class <- as.numeric(aero_aug$.pred_class)
```

```
aero_aug %>%
  roc_curve(truth = high_impact, .pred_class, event_level="second") %>%
  autoplot()
```



```
aero_aug %>%
  roc_auc(truth = high_impact, .pred_class, event_level="second")
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc binary      0.817
```

The first logistic regression is a good predictive model. The next step is to select fewer variables to see those increase the reliability of the model.

```
# Create recipe. Use all variables except aerobic_TE and related
aerobic_rec2 <- recipe(high_impact ~ middle_distance + avg_spd + min_elevation + avg_stride +
  avg_run_cadence + avg_hr, data = train_df2)

summary(aerobic_rec2)
```

```
## # A tibble: 7 x 4
##   variable      type    role    source
##   <chr>         <chr>  <chr>   <chr>
## 1 middle_distance nominal predictor original
## 2 avg_spd        numeric predictor original
## 3 min_elevation  numeric predictor original
## 4 avg_stride     numeric predictor original
```

```
## 5 avg_run_cadence numeric predictor original
## 6 avg_hr          numeric predictor original
## 7 high_impact     nominal outcome    original
```

```
log_reg <- logistic_reg() %>%
  set_engine("glm")
```

```
aero_wkfl2 <- workflow()%>%
  add_model(log_reg) %>%
  add_recipe(aerobic_rec2)
```

```
aero_fit2 <- aero_wkfl2 %>%
  fit(data = train_df2)
```

```
tidy(aero_fit2)
```

```
## # A tibble: 7 x 5
##   term                estimate std.error statistic  p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)      -133.      34.0      -3.93 8.63e- 5
## 2 middle_distanceY   2.72     0.358       7.60 3.07e-14
## 3 avg_spd           -14.7     4.15      -3.55 3.85e- 4
## 4 min_elevation     -0.0101   0.00576   -1.75 7.94e- 2
## 5 avg_stride         77.6     23.3       3.33 8.57e- 4
## 6 avg_run_cadence    0.764    0.224       3.40 6.63e- 4
## 7 avg_hr             0.185    0.0444      4.17 2.99e- 5
```

```
predict(aero_fit2, test_df2)
```

```
## # A tibble: 107 x 1
##   .pred_class
##   <fct>
## 1 0
## 2 0
## 3 0
## 4 0
## 5 0
## 6 0
## 7 0
## 8 1
## 9 0
## 10 1
## # ... with 97 more rows
```

```
aero_aug2 <- augment(aero_fit2, test_df2)
```

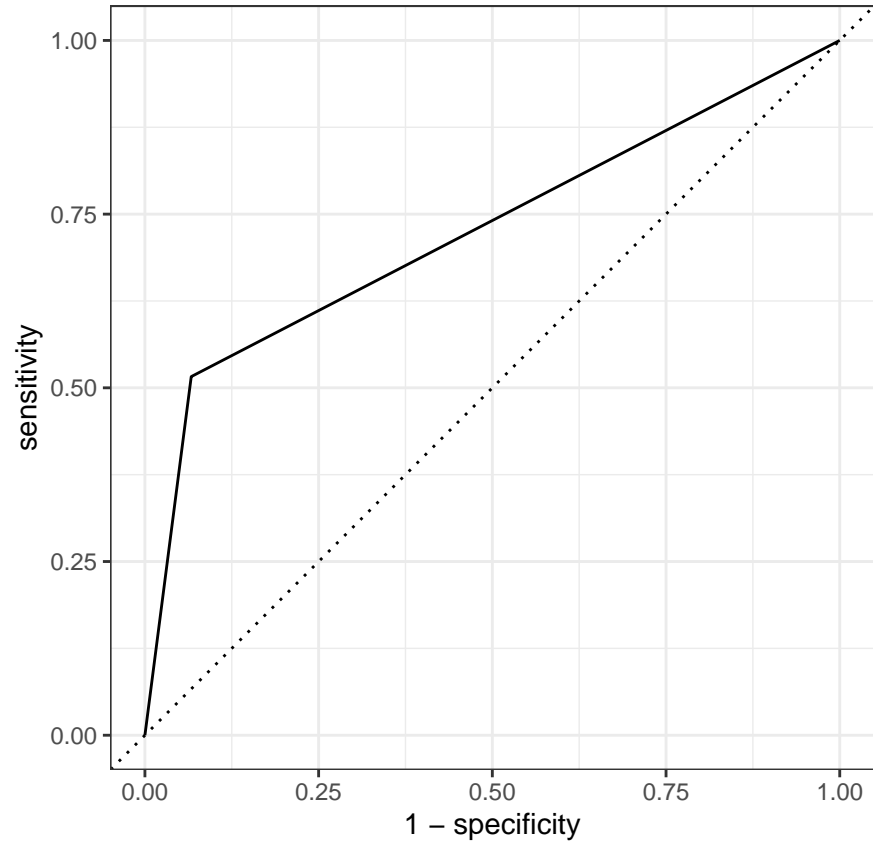
```
aero_aug2 %>% select(high_impact, .pred_class)
```

```
## # A tibble: 107 x 2
##   high_impact .pred_class
##   <fct>      <fct>
## 1 0          0
```

```
## 2 0      0
## 3 0      0
## 4 0      0
## 5 0      0
## 6 0      0
## 7 0      0
## 8 1      1
## 9 0      0
## 10 1     1
## # ... with 97 more rows
```

```
aero_aug2$.pred_class <- as.character(aero_aug2$.pred_class)
aero_aug2$.pred_class <- as.numeric(aero_aug2$.pred_class)

aero_aug2 %>%
  roc_curve(truth = high_impact, .pred_class, event_level = "second") %>%
  autoplot()
```



```
aero_aug2 %>%
  roc_auc(truth = high_impact, .pred_class, event_level = "second")
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc binary      0.725
```

Both of these models are acceptable for predicting whether a run highly impacts performance. These analyses provide a good starting point for building a more complex model that can predict good performance. The possible next step is to use k-fold cross validation to predict when good performance will happen given a series of events.